



NMRA Technical Note	
Layout Command Control™ (LCC) CAN Frame Transfer	
July 22, 2024	TN-9.7.2.1

Adopted as a NMRA Technical Note

The OpenLCB Standard document appended to this cover sheet has been formally adopted as a NMRA Standard by the NMRA Board of Directors on the date shown in the *Adopted* column in the *Version History* table below.

Version History

Date	Adopted	Summary of Changes
Feb 17, 2015		Initial version submitted for public comment
Feb 6, 2016	Feb 20, 2016	Minor grammatical corrections and readability improvements
Apr 25, 2021	July 2, 2021	Changed LCC logo to include the ® symbol Changed “Layout Command Control” to have the ™ symbol Added the NMRA Legal Disclaimer fine-print Changed the OpenLCB license to “Creative Commons Attribution-ShareAlike 4.0 International”
July 22, 2024		Add to the behavior of the AME Message

Important Notices and Disclaimers Concerning NMRA Standards Documents

The Standards (S), Recommended Practices (RP), Technical Note (TN) and Technical Information (TI) documents of the National Model Railroad Association ("NMRA Standards documents") are made available for use subject to important notices and legal disclaimers. These notices and disclaimers, or a reference to this page, appear in all standards and may be found under the heading "Important Notices and Disclaimers Concerning NMRA Standards Documents."

Notice and Disclaimer of Liability Concerning the Use of NMRA Standards Documents

NMRA Standards documents are developed within the Standards and Conformance Department of the NMRA in association with certain Working Groups, members, and representatives of manufacturers and sellers. NMRA develops its standards through a consensus development process, which brings together volunteers representing varied viewpoints and interests to achieve the final product. NMRA Standards documents are developed by volunteers with modeling, railroading, engineering, and industry-based expertise. Volunteers are not necessarily members of NMRA, and participate without compensation from NMRA.

NMRA does not warrant or represent the accuracy or completeness of the material contained in NMRA Standards documents, and expressly disclaims all warranties (express, implied and statutory) not included in this or any other document relating to the standard or recommended practice, including, but not limited to, the warranties of: merchantability; fitness for a particular purpose; non-infringement; and quality, accuracy, effectiveness, currency, or completeness of material. In addition, NMRA disclaims any and all conditions relating to results and workmanlike effort. In addition, NMRA does not warrant or represent that the use of the material contained in NMRA Standards documents is free from patent infringement. NMRA Standards documents are supplied "AS IS" and "WITH ALL FAULTS."

Use of NMRA Standards documents is wholly voluntary. The existence of an NMRA Standard or Recommended Practice does not imply that there are no other ways to produce, test, measure, purchase, market, or provide other goods and services related to the scope of the NMRA Standards documents. Furthermore, the viewpoint expressed at the time that NMRA approves or issues a Standard or Recommended Practice is subject to change brought about through developments in the state of the art and comments received from users of NMRA Standards documents.

In publishing and making its standards available, NMRA is not suggesting or rendering professional or other services for, or on behalf of, any person or entity, nor is NMRA undertaking to perform any duty owed by any other person or entity to another. Any person utilizing any NMRA Standards document, should rely upon their own independent judgment in the exercise of reasonable care in any given circumstances or, as appropriate, seek the advice of a competent professional in determining the appropriateness of a given NMRA Standards documents.

IN NO EVENT SHALL NMRA BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO: THE NEED TO PROCURE SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE PUBLICATION, USE OF, OR RELIANCE UPON ANY STANDARD OR RECOMMENDED PRACTICE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE AND REGARDLESS OF WHETHER SUCH DAMAGE WAS FORESEEABLE.

Translations

NMRA's development of NMRA Standards documents involves the review of documents in English only. In the event that an NMRA Standards document is translated, only the English version published by NMRA is the approved NMRA Standards document.

Official Statements

A statement, written or oral, that is not processed in accordance with NMRA policies for distribution of NMRA communications, or approved by the Board of Directors, an officer or committee chairperson, shall not be considered or inferred to be the official position of NMRA or any of its committees and shall not be considered to be, nor be relied upon as, a formal position of NMRA.

Comments on Standards

Comments for revision of NMRA Standards documents are welcome from any interested party, regardless of membership. However, **NMRA does not provide interpretations, consulting information, or advice pertaining to NMRA Standards documents.**

Suggestions for changes in documents should be in the form of a proposed change of text, together with appropriate supporting comments. Since NMRA standards represent a consensus of concerned interests, it is important that any responses to comments and questions also receive the concurrence of a balance of interests. For this reason, NMRA, its departments, Working Groups or committees cannot provide an instant response to comments, or questions except in those cases where the matter has previously been addressed. For the same reason, NMRA does not respond to interpretation requests. Any person who would like to participate in evaluating comments or in revisions to NMRA Standards documents may request participation in the relevant NMRA working group.

Laws & Regulations


Users of NMRA Standards documents should consult all applicable laws and regulations. Compliance with the provisions of any NMRA Standards document does not constitute compliance to any applicable regulatory requirements. Implementers of the standard are responsible for observing or referring to the applicable regulatory requirements. NMRA does not, by the publication of NMRA Standards documents, intend to urge action that is not in compliance with applicable laws, and NMRA Standards documents may not be construed as doing so.

Copyrights

NMRA Standards documents are copyrighted by NMRA under US and international copyright laws. They are made available by NMRA and are adopted for a wide variety of both public and private uses. These include both use, by reference, in laws and regulations, and use in private self-regulation, standardization, and the promotion of modeling, structural and engineering practices and methods. By making NMRA Standards documents available for use and adoption by public authorities and private users, NMRA does not waive any rights in copyright to the NMRA Standards documents.

IMPORTANT NOTICE

NMRA Standards documents do not guarantee or ensure safety, security, health, or environmental protection, or ensure against interference with or from other systems, devices or networks. NMRA Standards documents development activities consider research and information presented to the standards development group in developing any safety recommendations. Other information about safety practices, changes in technology or technology implementation, or impact by peripheral systems also may be pertinent to safety considerations during implementation of the standard. Implementers and users of NMRA Standards documents are responsible for determining and complying with all appropriate safety, security, environmental, health, and interference protection practices and all applicable laws and regulations.



OpenLCB Technical Note

CAN Frame Transfer

July 22, 2024

Adopted

1 Introduction

This Technical Note contains informative discussion and background for the corresponding “OpenLCB CAN Frame Transfer Standard”. This Technical Note is not normative in any way.

For information on format and presentation conventions, see:

- 5 • OpenLCB Common Information Technical Note

The Frame Transfer layer of the OpenLCB-CAN stack lives above the CAN Physical Layer, and below the Message Network layer. It is responsible for ensuring the reliable transport of the CAN frames that make up OpenLCB-CAN messages.

CAN provides reliable transport between any two nodes on a CAN segment within limitations:

- 10 1. Every CAN header must be unique by construction. Collisions between frames with identical headers and no data can result in lost frames; collisions between frames with identical headers that do contain data do not result in well-defined behavior. Both cases must be avoided by design to avoid intermittent problems.
- 15 2. CAN does not provide any indicator of which node sent a particular frame. Any node can send any frame, and the receiving nodes have no CAN-provided mechanism for identifying the source.
3. It's useful to have a way of addressing a frame to be processed by a specific node, but all CAN nodes receive every frame.
- 20 4. CAN does not provide a “link up” or “link down” notification. Nodes may come and go from a CAN segment at any time, and on an individual basis. In general, one can't assert that a particular node is always present, always comes up first, or never comes up first.
- 25 5. On a busy segment, CAN frames are sent in a strict priority order, with the priority determined by the content of the frame headers. If multiple nodes have frames to send at the same time, the highest priority (lowest numerical value header) frame will be sent first. There is no concept of “reserved bandwidth” except through frame priority, and low-priority frames may require significant time to be sent.

1.1 CAN Specialization for OpenLCB

CAN is adapted to OpenLCB use by a series of specifications which together mitigate the above limitations:

- 30 1. Putting a unique source ID field in each frame ensures that the frames are unique (item 1 above), and provides a frame-level indicator of which node sent the frame for the purposes of monitoring, debugging, etc. (item 2).
2. Providing a destination address that allows a standard method of addressed communications (item 3).
- 35 3. Using a distributed algorithm to determine the value for that ID field prevents needing a single “manager” to provide it, thereby handling item 4.
4. The header bit fields discussed below are organized to handle item 5.

40 Further, OpenLCB-CAN is required to work with standard CAN components (e.g., interface chips, bridges, computer adapters). These standard components must not require customization to operate with OpenLCB. This forces a stringent requirement on protocol design, in that OpenLCB-CAN cannot require specific timing, deliberate creation of error cases or non-standard error handling.

OpenLCB-CAN is required to work without user intervention; user presets are not required for the CAN frame communications layer to work. The Node ID that's needed for communications is required to be pre-configured into the node.

- 45 It's desirable for OpenLCB-CAN nodes to detect duplicate (non-unique) Node IDs, but 100% reliable detection is not required.

1.2 CAN collisions in loaded networks

50 Because of the CAN priority arbitration structure, the network can run at high utilization rates. Under those circumstances, all nodes attempt to send their highest-priority frame, and the node with the absolute highest priority frame (first dominant bit) wins. All other nodes wait until just after that frame and try again.

This behavior greatly increases the probability of collisions when multiple nodes attempt to send frames with the same header.

55 For example, consider 5 nodes each of which want to send the high-priority frames 1, ..., 5, plus one more node A that wants to send the low priority frame 15. If all those nodes attempt to send at once, the frames will arbitrate out and will be sent in the order 1, 2, 3, 4, 5, 15. If another node, B, also wants to send a frame 15, and starts to send it at any point during that entire sequence, it will end up colliding with the 15 that is being sent by A.

60 Particularly during layout startup, when many nodes have frames to send, this will result in an unacceptable probability of error, unless the protocols have been designed to have as few header collisions as possible, and absolutely no data collisions.

2 Annotations to the Standard

This section provides background information on corresponding sections of the Standard document. It's expected that two documents will be read together.

65 **2.1 Introduction**

(Nothing to add to the Standard)

2.2 Intended Use

70 The OpenLCB protocol suite uses 6-byte Node ID values, but this protocol layer could be useful in other contexts too. The Standard is therefore written so that any Node ID length through $7 \times 12 = 84$ bits can be used in a straight-forward manner.

2.3 References and Context

75 Note that the Standard references the Unique ID Standard, but only to the extent that each node is required to have a unique ID as its node ID. The Standard would work equally well with a Unique ID that was allocated via any mechanism, and with any length from 12 through 84 bits. OpenLCB uses a 48-bit unique ID as a node ID, which is assumed for the rest of this Technical Note.

2.4 Frame Format

Standard header frames are also known as CAN 2.0A frames. Extended header frames are also known as CAN 2.0B frames.

80 OpenLCB uses a straight-forward transfer of extended-header frames, using no special features, to reduce complexity and increase robustness.

The Standard is silent on the uses for standard CAN (11-bit header) frames. The proper-operation requirement is so nodes will tolerate them in case they're needed for other purposes. The Atmel bootloader¹ uses standard frames, for example, and the Microchip bootloader² can also use them. Allowing standard frames then allows the use of these bootloaders in parallel with OpenLCB-CAN. 85 Similarly, CBUS³ uses standard frames, and therefore could coexist with OpenLCB on a CAN segment.

This proper-operation requirement could also have been phrased as “shall ignore ...”, but the current phrasing is thought to be more exact.

90 RTR frames are not used in the protocol because CAN semantics require a specific use for them, which has been built into some silicon implementations. There are also some arbitration issues, see: CiA Application Note 802 (2005)⁴ for more information.

A CAN node can emit an overload frame when “Due to internal conditions, the node is not yet able to begin reception of the next (frame). A node may generate a maximum of two sequential overload frames to delay the start of the next (frame).” Those frames are 17 to 23 extra bit times each.⁵ The

¹See Atmel application note “AVR076: AVR® CAN - 4K Boot Loader”

<https://www.microchip.com/content/dam/mchp/documents/OTH/ApplicationNotes/ApplicationNotes/doc8247.pdf>

²See Microchip application note “AN247 A CAN Bootloader for PIC18F CAN Microcontrollers”

<http://ww1.microchip.com/downloads/en/AppNotes/00247a.pdf>

³See <https://www.merg.org.uk/resources/cbus>

⁴See <http://www.can-cia.org/>

⁵See <http://rs232-rs485.blogspot.com/2009/11/can-bus-message-frames-overload.html>

95 protocol does not specify transmission of overload frames because not all CAN controller hardware can deliberately send them. We require that nodes be able to handle them because some CAN controller hardware will occasionally send them automatically.

100 People will occasionally refer to a restriction against having seven 1 bits in the top 11 bits of the CAN header. There is no such restriction, and the statements are the result of a misunderstanding. There is a restriction against having seven consecutive recessive bits on the CAN segment, but sending seven consecutive 1 bits will not result in this. CAN uses a bit-stuffing technique to prevent that by inserting a dominant bit on the line after the fifth consecutive recessive bit, and removing it at the receiver.

The CAN format has been allocated on nibble boundaries to make it easier to read dumps of packets. The header is considered to be right (LSB) aligned.

105 The reserved most-significant bit (MSB) is likely to be used as a priority-boost bit in the future, for example so that a gateway can gain priority access to the CAN segment for various operations that require synchronization or atomicity. CAN encodes “do first” priority as 0 and “do later” as 1, so the Standard requires that a 1 bit be sent. To preserve the utility of this, nodes must not require either a 1 or 0 on receipt. It is therefore not available as a protocol expansion bit.

110 When zero, the second-most-significant bit (0x0800,0000) indicates that the frame is for local control on the CAN segment. Those frames should never leave the CAN segment. When one, it indicates that the frame is carrying OpenLCB message traffic, which can leave the CAN segment to proceed to the rest of the network; see the message-level documentation for more information on that.

115 The order of the header fields was chosen to get proper priority and access disambiguation via CAN's standard mechanisms.

2.5 States

120 In the Inhibited state, a node can only communicate on the CAN segment to allocate its Node ID alias(es). This involves sending CID and RID frames with a tentative Node ID alias value. Once that process is complete, the node has an assigned Node ID alias and can transition to the Permitted state, where all communications are possible.

If a node fails, it should restart in the Inhibited state. In general, this state transition is not visible to other nodes until the failed node starts communicating.

125 Designers of nodes that maintain alias pools, for example in gateways, need to carefully consider to what extent the state applies to the entire node, and to what extent it applies to each individual alias in the gateway pool. Anything that transmits OpenLCB frames on CAN needs to have reserved an alias and have transitioned to Permitted state by transmitting an Alias Map Definition frame using that particular alias. It can acquire and stock-pile other aliases for later use, putting them into Permitted state either immediately or at some later time.

2.6 CAN-specific Control Frames and Interactions

130 2.6.1 Control Frame Format

RID, CID, etc, frames are assigned low numbers to give them higher transmission priority.

The coding is done to allow up to 7 separate CID sequence numbers, for use by protocols with longer node IDs.

135 The coding of the 0x07NN field was selected to allow expansion of control fields in the 0x070N range, and expansion of error indications in the 0x071N range.

The block of CAN frames with Content Field from 0x1000 to 0x3FFF are not used by OpenLCB. Although formally reserved, they can serve as a convenient 26-bit (with bit 28) address block for experimentation and migration of other equipment.

2.6.2 Interactions

140 Communications protocols are about more than formats. The protocols must also specify the interactions in which the frames are used.

2.6.2.1 Reserving a Node ID Alias

145 At the highest level, this algorithm is broadcasting the tentative alias to see if any other node is checking the same tentative alias. If it is in use, the algorithm tries additional ones until it finds one that's available. The following considerations had to be taken into account when designing the alias reservation algorithm.

150 CAN arbitration reliably avoids collisions between frames with unique headers. It does not guarantee non-overlapping transmission of frames from separate nodes with identical headers and no data; if the timing is right, they may overlay each other such that only one frame appears to have been sent. It is very difficult to ensure that two nodes always send initialization packets only at different times, particularly given the load-related frame synchronization issue described earlier. In addition, CAN will eventually signal an error if two packets with the same header and different data payloads collide, but not all CAN interface hardware provides reliable indications about why that error occurred.

155 If a node sends out a check frame containing just the alias, then it could expect that another node to complain if it has the same alias. This would work, except in the (admittedly unlikely) case where both nodes send out the identical check frame simultaneously. Neither would recognize a conflict and both would consider that they own the same alias. Therefore, a method needed to be found that guaranteed that the check frame(s) are unique, even if they are sent simultaneously.

160 The CID/RID algorithm for finding aliases is designed to work with CAN arbitration, but without causing any CAN errors. During execution of the algorithm, some frames might have the same header. They therefore have to have the same contents, including null contents, or else a non-arbitration error can occur.

165 If two nodes have taken the same tentative alias, at least one of the packets used in the broadcast will not be identical between the two nodes, because their Node IDs are different in at least one bit. CAN will successfully arbitrate this, and one node will receive the frame sent by the other node, causing it to back off.

170 CAN transmissions are not atomic operations; a node can receive a frame between the time it tells the hardware to send a frame and the time that frame is actually sent. It's therefore possible for both nodes contending for the same alias value to back off and try again. The mechanism for generating the next alias value to try must generate different values for different nodes to prevent this from becoming a permanent lock-step error.

The delay at the end of the algorithm is to ensure that higher latency nodes, such as software nodes working through USB converters, can reply to CID frames from nodes that come up on a working link.

175 OpenLCB is a soft-real-time system, and software that's interacting with it needs to have response times of a couple tens of milliseconds or better to be reliable. The algorithm provides a wait of a few times that to enable those programs to take part.

Nodes may assign different alias values each time they are powered up. There is no requirement that they always reserve and use the same one, nor is it possible to guarantee that they will.

180 The sequence number MMM is sent in descending order, 0x7, 0x6, ..., to give priority on the CAN segment to frames later in the process. This lets some nodes complete as early as possible and remove themselves from contention.

Although it might be a useful debugging tool, putting the full Node ID in the data portion of the frames can cause errors. Therefore, we decided not to do that.

185 Reserving a node ID alias value is separate from actually using it to take a node from Inhibited to Permitted state. A gateway device, for example, might want to stockpile a few reserved alias values so that they can be rapidly associated with a specific node that wants to communicate. The node ID that's broken into four parts for the CID values needs to be unique, but it does not have to be the one that's eventually attached to the alias via an Alias Map Definition frame later on.

For example, a minimal sequence is four CID frames and a RID:

190 0x17123FED
 0x16456FED
 0x15789FED
 0x14ABCFED
 0x10700FED

195 The red nibbles show that these are four CID and one RID frames. The blue nibbles are the source ID alias being tested and allocated. The yellow nibbles are the underlying Node ID of the node, in this case 12 34 56 78 9A BC.

2.6.2.2 Transition to Permitted State

200 Note that there is no requirement that the node's alias value be consistent from one power-up reset to the next.

2.6.2.3 Node ID Alias validation and invalidation

205 The Alias Mapping Enquiry provides a mechanism for determining which alias is associated with each Node ID. This can be used when a new node attaches and needs the information. Operating nodes can also track this information by watching CAN segment traffic for control frames, without having to use the Alias Mapping Enquiry protocol.

The AME has two distinct usages:

- AME without Node ID: all active nodes reply with AMD and their Node ID;

- AME with Node ID: the receiving node which has a matching Node ID replies with AMD and its Node ID;

210

Node ID alias validation provides a way for a node joining a segment to rapidly learn the Node IDs and Node ID aliases of the reachable nodes. This may be useful e.g. for a gateway that is connecting two segments that are already in operation, or a monitor node that wants to determine that all nodes have come up properly.

- 215 In general, simple or regular nodes won't need to do node alias validation. Those nodes don't need to keep track of the Node ID to alias mapping, because they simply respond to frames they receive that contain an alias. They just work with the alias itself. They also don't need to track other nodes' state, etc.

- 220 In order to acquire the map between CAN aliases and node IDs, a gateway needs to be able to send a CAN frame requiring "everybody reply with your node ID". Much like Ethernet gateway mapping, this needs to return the node ID of just the specific CAN-attached hardware, not all the node IDs that can be reached through one or more gateways, so it needs to be a segment-specific frame defined only at the CAN level. It must not be a OpenLCB common message, though higher level protocols may want to provide a similar capability that spans the entire OpenLCB installation.

- 225 In either Inhibited or Permitted state, receipt of an Alias Mapping Enquiry frame with no data content invalidates any previously validated alias to/from Node ID mapping information. This will be followed by nodes in Permitted state responding with AMD frames containing their alias and Node ID, which will allow a node to rebuild the valid entries any alias ↔ Node ID mapping it maintains. Nodes which have transitioned to Inhibited state or left the network will not reply. Sending an AME with no data
230 content thus ensures that nodes are using only currently accurate alias information.

- For an example of how this is useful, consider virtual nodes that are created and maintained by some host node, e.g. in the Train Control protocol. Other nodes are referencing those virtual nodes through aliases. When the control node is rebooted, it no longer knows that those virtual nodes have been created, nor does it know that those aliases refer to them, yet the other nodes are still sending messages
235 using those aliases. Those messages don't get delivered to the virtual nodes because those virtual nodes don't exist yet. By sending an AME with no content when it reboots, the host node can ensure that the aliases for those virtual nodes are invalidated in the other nodes: The virtual nodes won't be replying to the AME with no content because they don't exist yet. When another node next tries to send a message to a virtual node, it will realize that it doesn't know the alias of that virtual node. It will
240 go through the process of using an AME to acquire that alias, which will prompt recreation of the virtual node so that operations can continue.

2.6.2.4 Transition to Inhibited State

The Alias Map Reset frame indicates that the alias value in the source address field may no longer be associated with the same Node ID, and mapping tables should be reset.

- 245 The full node ID in the Alias Map Reset frame allows nodes that only are aware of an alias for their conversation partner to convert that alias to a full node ID for later use.

This is a way to release or reuse an alias, while informing everybody that the node is no longer reachable.

250 The alias is only released if this node allows some other node to allocate it via CID/RID frames. The current node can still hang on to it if desired, for example to reuse it with another Node ID, by replying to CID/RID frames from another node attempting to reserve it.

255 Gateways have to obtain unique aliases for remote nodes they are proxying onto the segment. Gateways maintain the mapping between remote node's node ID and the local alias. If they need to break that mapping, e.g. because they need to re-purpose a local alias due to resource limitations in the mapping tables, they must send a "Alias Map Reset" frame to force nodes to drop their alias information. "Alias Map Reset" is a CAN-specific frame limited to the segment, but all gateways on the segment must act on it. This is then followed by an "Alias Map Definition" when the alias is allocated to a new node ID.

260 Regular nodes won't need to do this. They stay active until they go instantly offline due to e.g. power-off.

265 When a node goes to inhibited state or releases an alias by sending an Alias Map Reset frame, that's telling other nodes that it can no longer be reached by using that alias. The last paragraph in this section of the Standard gives the node receiving the Alias Map Reset 100msec to finish processing and stop using the alias. If it still needs to send messages to the node that sent the Alias Map Reset, the receiving node can attempt to reacquire a new alias via an Alias Map Enquire frame or higher-level message(s). If the alias has changed, that will work, but if the sending node has transitioned to Inhibited state there will be no reply to the attempt to obtain the new alias.

2.6.2.5 Node ID Alias Collision Handling

270 During normal operation in the Permitted state on a single CAN segment, a node should never encounter a Node ID alias collision.

They can happen when network topology changes, for example when connecting two operating CAN segments. In that case, nodes on the two segments have assigned aliases independently, and can have used the same value. The underlying Node IDs do not collide, because they are assigned globally uniquely, but that's not sufficient to ensure that the shorter aliases are independent.

275 In the case of a collision, one or both nodes will release the current alias. If it's the alias that was used to go to Permitted state, e.g. the current alias being used by this node, the node must also drop back to Inhibited state. The node(s) then can, but is not required to, return to Permitted state with a new alias. In the process the node(s) will emit "Alias Map Reset" and "Alias Map Definition" frame that will allow any gateway(s) to update their maps.

280 The preceding paragraph and the language of the Standard is complicated because it has to cover the case of a gateway or similar node that has reserved a number of node ID aliases, but is not currently using them as a node. In other words, the reservation CID/RID process has completed to reserve the alias, but no mapping of the alias to a specific node ID has been published.

285 Simple leaf nodes don't keep internal mappings between full node IDs and aliases, so this transition will not cause any problems for them unless they are currently in communication with another node,

e.g. in the middle of a sequence of frames for a datagram or stream. In that case, the node may have to invoke higher-level error handling to cause the sequence to be retried.

290 For example, a detachable throttle will detect that it's reattaching, reserve an alias value, and then enter permitted state with that value. Since the reservation happened while connected to the current CAN segment, there is no possibility of a duplicate.

Example: Nodes A and B are in Permitted mode with different node IDs, but the same alias values. Some gateway node C requests a list of all nodes on the CAN segment:

C: 0x10702ccc (AME frame)

A: 0x10701nnn 01 02 03 04 05 06 (AMD frame with Node ID 0x0102,0304,0506)

295 (For simplicity, B has not transmitted yet) B discovers that it's alias has been duplicated.

B: 0x10703nnn 55 AA 44 AA 33 AA (AMR frame with Node ID 0x55AA,44AA,33AA)

Node A also discovers that it's alias has been duplicated.

A: 0x10703nnn 01 02 03 04 05 06 (AMR frame with Node ID 0x0102,0304,0506)

B then goes back into operation by reserving a new alias value and transitioning to Permitted state:

300 B: 0x1755Ammm (1st CID)

B: 0x16A44mmm (2nd CID)

B: 0x15AA3mmm (3rd CID)

B: 0x143Ammm (4th CID)

Node A will also reserve a new alias.

305 A: 0x17010kkk (1st CID)

A: 0x16203kkk (2nd CID)

A: 0x15040kkk (3rd CID)

A: 0x14506kkk (4th CID)

After 200 ms both nodes finish their alias reservation and transition to Permitted state.

310 B: 0x10700mmm (RID)

B: 0x10701mmm 55 AA 44 AA 33 AA (AMD with B's node ID attached)

A: 0x10700kkk (RID)

A: 0x10701kkk 01 02 03 04 05 06 (AMD with A's node ID attached)

2.6.2.6 Duplicate Node ID Handling

- 315 This should never happen, of course, but it's possible that two nodes will have the same Node ID assigned. This will result in frame loss and operational difficulties. OpenLCB does not guarantee prompt detection of this condition, at any level of the protocol stack, but it's extremely likely that it will be detected. At the Frame Transfer level, two nodes on the same CAN segment sharing a common Node ID will be detected when a node receives a Alias Map Definition frame that contains its own
- 320 Node ID.

The “in addition to any other actions” refers to e.g. acting on identical alias values, should that also be the case.

- 325 Sending the PCER message is level jumping, in that such messages are defined at a much higher level in the protocol stack. We include it here because it allows a much cleaner set of error handling operations using the event producer/consumer techniques associated with OpenLCB. This decision can be revisited.

After sending the (optional) PCER message, the node has to go offline so that it doesn't interact anymore until the user resets it to restart operation. This is necessary to prevent the PCER itself triggering additional messages back, resulting in a never-ending cycle.

- 330 Example: Nodes A and B are in Permitted mode with the same node IDs, but different aliases. Some gateway node C requests a list of all nodes on the CAN segment:

C: 0x10702ccc (AME frame)

A: 0x10701aaa 01 02 03 04 05 06 (AMD frame with Node ID 0x010203040506)

(For simplicity, B has not transmitted yet) B discovers that its node ID has been duplicated.

- 335 B: 0x195B4bbb 01 01 00 00 00 00 02 01 (Producer-Consumer Event Report (PCER) message with event ID 0x0101,0000,0000,0201)

If the two nodes have also assigned the same alias value, that interaction is also invoked here.

2.6.2.7 Reporting CAN Link Status

- 340 The CAN Standard defines the “Error Active”, “Error Inactive” and “bus off” states, along with extensive fault-confinement rules. By using qualified CAN implementations, OpenLCB-CAN nodes pick these behaviors up. The interactions in this section are then about reporting the results of the CAN Standard operations. The “Error Information Report 0” through “Error Information Report 3” control frames indicate specific transitions in the node's state. No OpenLCB-CAN node is required to react to those reports. The reports only provide information.

- 345 Designers of monitoring equipment that attaches to the CAN link can use this information directly. Designers of gateway nodes that connect CAN segments together or connect CAN segments to other types of OpenLCB links may want to convert these control frames (which are not allowed to go off their original CAN segment) to higher level messages. There are well-known EventIDs for this purpose, but that use is beyond the scope of this Standard and Technical Note.

350 Since CAN nodes are required by the CAN standard to initialize in Bus Active state, no Error Information N frame(s) can be emitted at startup. They cannot be emitted until there's an actual state transition in the interface.

It's not possible for a node to emit a frame to indicate that it has entered the "Bus Off" state, so there's no message corresponding to that.

355 **2.6.3 Node ID Alias Generation**

We have a preferred implementation for node ID alias generation which is described in Section 6 "Preferred Alias Allocation Method" below.

The Standard is written to specify the minimum needed to ensure proper operation. This allows developers to create other implementations if needed.

360 The term "of the same type" is meant to cover "from the same manufacturer". The idea is that manufacturer will likely number their boards sequentially, and it's possible that an OpenLCB installation will have multiple boards from the same manufacturer. This rule just says that the low bits of the sequential number appear in the alias, so that boards from a single manufacturer are not likely to collide.

365 Two other criteria that were considered:

- The sequence of alias values generated by a node shall depend on the node's node ID in such a way that the sequence differs from the sequence generated by every other node.
- Given sufficient iterations, the sequence of alias values generated by a node shall include every valid alias value.

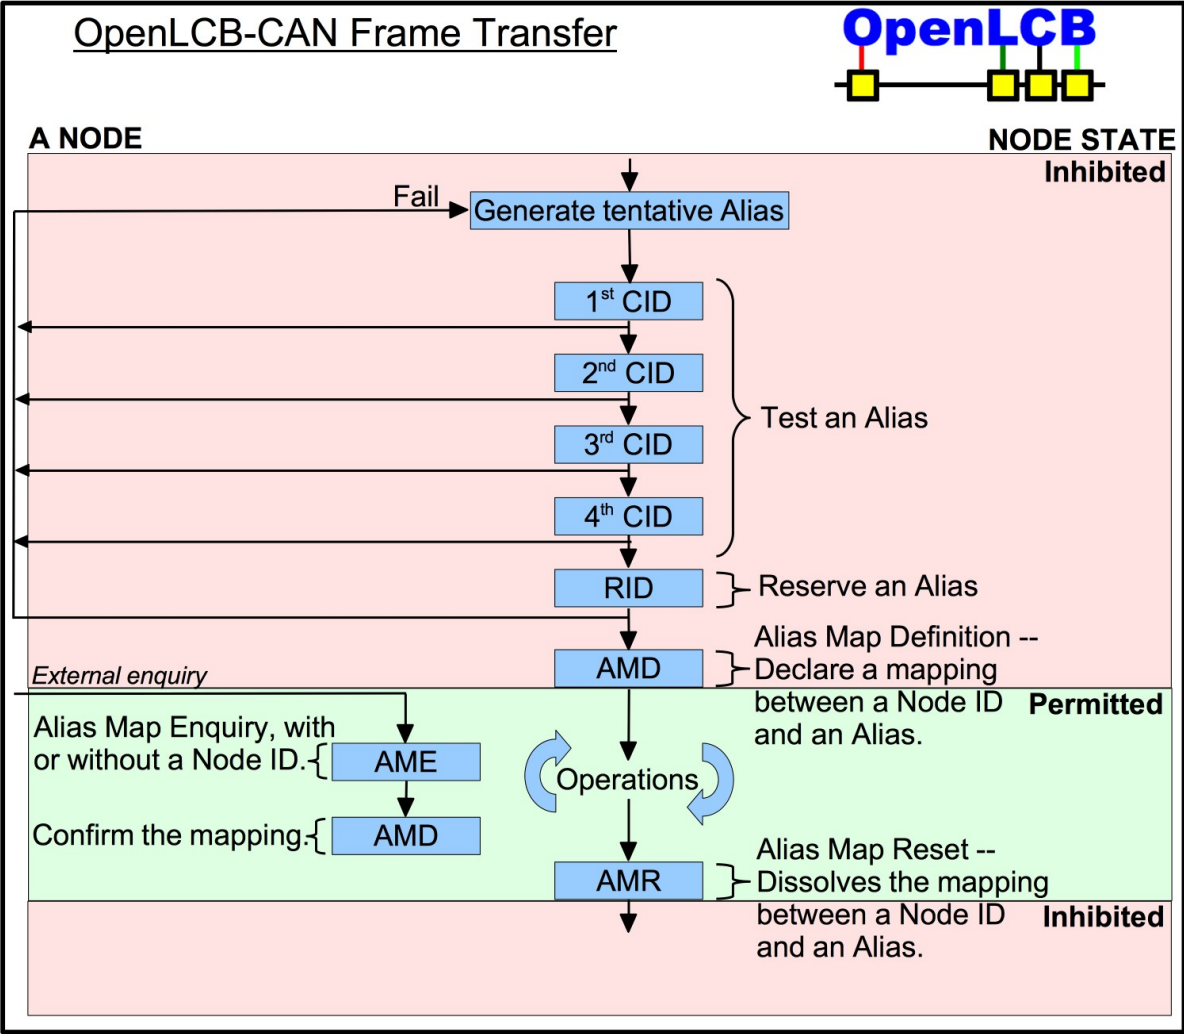
370 Although useful criteria, they were omitted as too hard to comprehensively test.

What's really necessary, but hard to describe in standard language, is that two separate nodes need to have different sequences of generated numbers. If node A will try 3, 4, 5, 6, ... then it's important that node B try some other sequence. If B also generates values in the sequence 3, 4, 5, 6, ... then the two nodes can get locked to each other trying the same value, both rejecting it, and then trying the same new value.

Formally, if A_i is defined as a particular value in node A's sequence, B_i is defined as a particular value in B's sequence, then "if $A_i = B_i$, then $A_{i+1} \neq B_{i+1}$ most of the time" is the important condition.

It's not required that two sequential aliases be different. It's possible that the sequence generator will come up with the same value occasionally, and the RID/CID mechanism is constructed to handle that.

380 The sequence on the CAN bus will be clearer, and in some cases the allocation process will be faster, if nodes identify duplicate aliases in sequence, and just go on to the next one.



Note that in the diagram above, the 'Fails' are due to the node receiving a message which uses its tentative Alias, and it therefore needs to back off and try a new tentative Alias.

385

3 Gateways and segment management

Multiple CAN electrical regions can be connected into a single CAN segment via repeaters that allow the entire segment to share arbitration. The OpenLCB-CAN frame-level protocol was designed to work with these: It doesn't use RTR frames, and negotiates aliases over top of the CAN arbitration.

390 Bridges, which work by passing frames between separate segments, are also supported. In normal operation, OpenLCB-CAN frames can be transported through those bridges unchanged. While the

bridge is active, the alias negotiation system will work properly. If a bridge is inactive, e.g. the two segments are separate, while nodes are coming up and allocating their aliases, it's possible that duplicate aliases will be allocated. Bringing up the CAN bridge hardware before bringing up the nodes prevents duplicate assignment. If duplicates are assigned, e.g. due to bringing up a bridge between active CAN segments, the nodes will discover the duplicates on their first transmission, and will rectify the situation using the procedure defined in the "Node ID Alias Collision Handling" section.

Higher-capability bridges, and gateways to non-CAN OpenLCB networks, will want to handle the assignment of local aliases directly. This prevents alias collisions between separate CAN segments, and allows very large OpenLCB networks, but requires more work from the gateway. It will allocate a pool of aliases using the conventional technique, and then assign them to off-segment nodes as they need to communicate on the segment. Many algorithms are possible for that assignment.

If/when a full Node ID is needed, it can be obtained by sending an "Alias Mapping Enquiry" frame with an appropriate Destination Node ID in local alias form. The reply will eventually come back with the Source ID in local alias form and carrying the full Node ID in the frame body.

4 Throughput and Temporarily Deaf Nodes

29-bit extended header and 8-byte payload results in a total transmission taking about 135 bit times, varying slightly with bit-stuffing.

125Kbps with max-length extended header frames is about 900 frames/second. CAN is very good at running at 100% utilization, so long as nodes can keep up and a proper set of priorities is in place. The CAN network itself does not require any inter-frame spacing, limitations on bandwidth usage, etc.; the segment can be allowed to run at 100%.

OpenLCB requires that CAN-attached nodes be able to handle the full frame rate on the CAN bus. There is no guarantee that frames for a given node will arrive with non-zero time between them.

Some nodes are not always able to process frames. For example, the node may cease processing for a short time while writing non-volatile memory. Higher-level protocols, e.g. configuration write datagrams, have mechanisms built in to prevent overrun while writing configuration memory, but there could be multiple activities going on in parallel that will result in frames arriving while the node is not processing.

Short outages can be covered by CAN hardware buffering, so long as the node will eventually catch up even at full arrival rate. Outages long enough that frames are lost due to e.g. buffer overflow require a node to broadcast that it's back up if the hardware detected frames lost. This is because there could have been frames of some form that modified the OpenLCB node state, which are now lost (e.g. drop alias, CID/RID in absence, or even higher level events)

The real limitation is whether the nodes themselves can handle the full rate of frame arrival. There is (currently) no mechanism in OpenLCB to throttle the rate of frames arriving at a node, nor is it easy to create one in a CAN network. The stream and datagram protocol have been designed to allow a node to efficiently use a very limited amount of buffer memory, but don't provide mechanisms to limit the arrival of frames in general. Nodes must handle full rate frames, specifically including frames not addressed to them. For long term reliability, a node must be able to completely process an entire CAN

frame in the time it takes to receive the next one, which may be as little as 64 bit times, or about 500 microseconds for a header only frame.

435 Eventually OpenLCB may have to discuss use of Overload Frames to throttle for e.g. NVRAM writing, but various CAN hardware implementations may also be doing this already. This mechanism provides only a little more time, due to restrictions in the CAN specification that have become limitations on the silicon implementations.

5 On the choice of a 12-bit alias length

The size of the 12-bit NodeID alias field was selected by balancing two constraints:

- Smaller (fewer bits) allows more payload and/or simpler coding of other parts of messages.
- 440 • Larger allows more unique nodes to be accessed.

5.1 Issues

5.1.1 Addressing

445 The alias size limits the number of unique nodes that can take part in communications on the CAN segment. Because Node ID aliases are assigned independently on each CAN segment, the only issue is how many different nodes are involved, not which ones they are or what pattern(s) are available in their address numbers.

450 For electrical/timing reasons, a segment itself can only support a maximum of about 100 nodes⁶, but communication also involves nodes off the segment. For example, chaining N CAN segments via bridges will generally make up to 100*N nodes available. We have a use case of chaining small numbers of CAN segments together, implying the need to address 500 or 1000 nodes.

As another example, a very large network might connect many CAN segments via TCP/IP or other networked connections. Each CAN-attached node might need to communicate with a larger number of others spread across that network.

455 Finally, one can imagine connecting a few CAN segments in a star pattern to a central OpenLCB-CAN “backbone”, which must route the full cross-segment traffic.

In some cases, we want two aliases (source and destination), plus a few more bits, in the 29 bits total available in the header. Using 12-bits for each address leaves $29 - (2 * 12) = 5$ bits for these purposes, which seemed about right, and allows a maximum of 2^5 , or 32 nodes.

5.1.2 Collisions during Allocation

460 The alias allocation process resolves collisions (attempt to use an already allocated Node ID alias), but that takes time. Minimizing the number of collisions is good. If the address space is just the size of the number of things you want to address, the collisions get hard /slow to resolve across multiple nodes. This points to having a factor of 2, at least, between the number of nodes to be addressed and the size of the address space.

⁶The OpenLCB-CAN Physical Layer Standard says in its Intended Use section uses 50 nodes as an expected size. This is merely an informative statement of expectations, however, and some layouts will no doubt exceed that. We use the larger number 100 here to avoid placing any stronger limitation on what can be achieved electrically.

465 **5.1.3 Size**

Streaming provides the most stringent test of alias sizes. We'd like to have all 8 bytes of CAN frame payload available for transporting stream data, which requires putting both source and destination addresses, plus whatever control information is needed, into the header.

470 To get the full payload (8 bytes) for streaming requires getting the rest of the protocol control, including two aliases for source and destination, in the 29 bit header. Allowing 1 bit for priority/extension, 1 for protocol control, 3 for MTI and stream control (all very bare minimum) leads to a 12-bit node ID alias length.

6 Preferred Alias Allocation Method

This section documents the preferred algorithm for generating unique Node ID aliases.

475 CAN transmissions are not atomic operations; you can receive a frame between the time you tell the hardware to send a frame and the time that frame is actually sent. It's therefore possible for both nodes contending for the same alias value to back off and try again. Using the pseudo-random number generator as a sequence number makes it likely that the next attempt will be made using different alias values in the two nodes. The use of a sequence generator also makes the arbitration process faster in the case of sequential Node ID values. People like to use small and sequential numbers for things, and the sequence generator maps those to very different values that are less likely to collide during arbitration.

6.1 Preferred alias generation algorithm

The 12 bit alias has to be derived from a 48-bit sequence number in a way that keeps as much entropy (distinctiveness; randomness) as possible.

485 The entropy starts in the node ID values, which cover a 48-bit space (albeit non-uniformly). This provides an initial 48-bit value that can be mapped to a 12-bit value.

To create a sequence of values, you want a computationally cheap way of sequencing through all those possible 48-bit values. (Running through the 48-bit values and compressing each to a 12-bit alias value is less likely to get multiple nodes attempting the same sequence than just working through a 12-bit sequence) A suitable pseudo-random number generator does exactly this: it runs through full 48-bit sequence, with each node starting at separate point in the sequence, so the 48 bit values are unique.

A computationally convenient way to create a 12-bit value, while preserving as much of the randomness as possible, is to break the 48-bit value into four 12-bit values, align these, and exclusive-or them. Each bit-change in the original 48-bit number then becomes a change in the final 12-bit value.

495 Sequentially numbered nodes will have initial seeds that differ only in the low bits of the PRNG output. For shift-register PRNGs, this can present a problem, as all bytes of the result are shifting together with just one new bit shifting in at the top. (Traditionally, shift-register PRNGs shift toward the LSB) Only one bit changes. Combining bytes (or even 12-bit chunks) via XOR or add operations results in reduced entropy.

500 The OpenLCB proposed PRNG (next section) uses add operations with a constant term to avoid this problem, so treating the 48-bit PRNG output as four 12-bit values, which are then XORed, works well.

6.1.1 Preferred PRNG

The proposed 48-bit PRNG is from “A 48-Bit Pseudo-Random Number Generator”, Heidi G. Kuehn, Communications of the ACM, Volume 4, Issue 8 (August 1961), pages 350-352.

$$505 \quad x_{i+1} = (2^9+1) x_i + c$$

where $c = 29,741,096,258,473$ or $0x1B0CA37A4BA9$. The paper actually describes a PRNG that uses signed arithmetic, but for our purposes the 48th “sign” bit isn't important. This is a byte+bit shift and add, so it's quick to calculate.

Note that, unlike some other PNRGs, this one can generate a zero result, and can accept a zero seed.

510 It is initialized with the node ID as a 48-bit key, with the most significant byte of the node ID as the most-significant byte of the PRNG value.

6.1.2 Preferred PRNG to node ID alias mapping

The 12-bit node ID alias is made from the 48-bit PRNG value by splitting the PRNG value into four 12-bit values and XORing them together. This preserves the information content of every bit in the final result.

515

6.1.3 Example C implementation

The PRNG state is stored in two 32-bit quantities:

```
uint32_t lfsr1, lfsr2; // sequence value: lfsr1 is upper 24 bits, lfsr2 lower
```

The 6-byte unique Node ID is stored in the `nid[0:5]` array.

520 To load the PRNG from the Node ID:

```
lfsr1 = (nid[0] << 16) | (nid[1] << 8) | nid[2];
lfsr2 = (nid[3] << 16) | (nid[4] << 8) | nid[5];
```

To step the PRNG:

```
525 // First, form 2^9*val
uint32_t temp1 = ((lfsr1<<9) | ((lfsr2>>15)&0x1FF)) & 0xFFFFFFFF;
uint32_t temp2 = (lfsr2<<9) & 0xFFFFFFFF;
```

```
530 // add
lfsr2 = lfsr2 + temp2 + 0x7A4BA9L;
lfsr1 = lfsr1 + temp1 + 0x1B0CA3L;
```

```
// carry
lfsr1 = (lfsr1 & 0xFFFFFFFF) + ((lfsr2&0xFF000000) >> 24);
lfsr2 = lfsr2 & 0xFFFFFFFF;
```

535 Form a 12-bit alias from the PRNG state:

```
uint16_t LinkControl::getAlias() {
    return (lfsr1 ^ lfsr2 ^ (lfsr1>>12) ^ (lfsr2>>12) )&0xFFF;
}
```

7 Additional References

- 540 The OpenLCB-CAN Physical Layer TN contains useful references on the characteristics of CAN.
- B. Gaujal and N. Navet, “Fault confinement mechanisms on CAN: analysis and improvements”, IEEE Transactions on Vehicular Technology, pp.1103-1113, 54(3), May 2005 (An interesting analysis of how CAN nodes react to error rates on the CAN bus by e.g. taking themselves offline, etc.)
- 545 “A 48-Bit Pseudo-Random Number Generator”, Heidi G. Kuehn, Communications of the ACM, Volume 4, Issue 8 (August 1961), pages 350-352.

Appendix A: Calculation Examples

Starting 48-bit seed	Alias	Next 48-bit seed
00 00 00 00 00 00	0	1B 0C A3 7A 4B A9
1B 0C A3 7A 4B A9	11E	4F 60 3B 8B E9 52
4F 60 3B 8B E9 52	521	2A E3 F6 D8 D8 FB
2A E3 F6 D8 D8 FB	42E	0D DE 4C 05 1A A4
0D DE 4C 05 1A A4	464	E5 82 F9 B4 AE 4D

Table: Sequence of aliases and seeds starting from zero

Starting 48-bit seed	Alias	Next 48-bit seed	Next Alias
02 01 21 00 00 12	113	1F 4F C4 7A 6F BB	62D
02 01 12 00 00 21	113	1F 31 B5 7A 8D CA	A24
02 01 11 00 00 22	113	1F 2F B4 7A 8F CB	625
02 01 22 00 00 11	113	1F 51 C5 7A 6D BA	A2C

Table: Four seeds (of 2³⁶) that give the same alias, and their different subsequent seeds and aliases

Table of Contents

1 Introduction.....	1
1.1 CAN Specialization for OpenLCB.....	1
1.2 CAN collisions in loaded networks.....	2
2 Annotations to the Standard.....	2
2.1 Introduction.....	2
2.2 Intended Use.....	3
2.3 References and Context.....	3
2.4 Frame Format.....	3
2.5 States.....	4
2.6 CAN-specific Control Frames and Interactions.....	4
2.6.1 Control Frame Format.....	4
2.6.2 Interactions.....	5
2.6.2.1 Reserving a Node ID Alias.....	5
2.6.2.2 Transition to Permitted State.....	6
2.6.2.3 Node ID Alias validation.....	6
2.6.2.4 Transition to Inhibited State.....	7
2.6.2.5 Node ID Alias Collision Handling.....	8
2.6.2.6 Duplicate Node ID Handling.....	9
2.6.2.7 Reporting CAN Link Status.....	10
2.6.3 Node ID Alias Generation.....	10
3 Gateways and segment management.....	12
4 Throughput and Temporarily Deaf Nodes.....	12
5 On the choice of a 12-bit alias length.....	13
5.1 Issues.....	13
5.1.1 Addressing.....	13
5.1.2 Collisions during Allocation.....	14
5.1.3 Size.....	14
6 Preferred Alias Allocation Method.....	14
6.1 Preferred alias generation algorithm.....	14
6.1.1 Preferred PRNG.....	15
6.1.2 Preferred PRNG to node ID alias mapping.....	15
6.1.3 Example C implementation.....	15
7 Additional References.....	16