



NMRA Technical Note	
Layout Command Control™ (LCC) Simple Node Information	
July 22, 2024	TN-9.7.4.3

Adopted as a NMRA Technical Note

The OpenLCB Standard document appended to this cover sheet has been formally adopted as a NMRA Standard by the NMRA Board of Directors on the date shown in the *Adopted* column in the *Version History* table below.

Version History

Date	Adopted	Summary of Changes
Feb 17, 2015		Initial version submitted for public comment
Feb 6, 2016	Feb 20, 2016	Minor grammatical corrections and readability improvements Changed document numbering from TN-9.7.3.4 (Network & Transport Layer) to TN-9.7.4.3 (Session Layer)
Apr 25, 2021	July 2, 2021	Changed LCC logo to include the ® symbol Changed “Layout Command Control” to have the ™ symbol Added the NMRA Legal Disclaimer fine-print Changed the OpenLCB license to “Creative Commons Attribution-ShareAlike 4.0 International”
July 22, 2024		Updated how the two data blocks are versioned

Important Notices and Disclaimers Concerning NMRA Standards Documents

The Standards (S), Recommended Practices (RP), Technical Note (TN) and Technical Information (TI) documents of the National Model Railroad Association ("NMRA Standards documents") are made available for use subject to important notices and legal disclaimers. These notices and disclaimers, or a reference to this page, appear in all standards and may be found under the heading "Important Notices and Disclaimers Concerning NMRA Standards Documents."

Notice and Disclaimer of Liability Concerning the Use of NMRA Standards Documents

NMRA Standards documents are developed within the Standards and Conformance Department of the NMRA in association with certain Working Groups, members, and representatives of manufacturers and sellers. NMRA develops its standards through a consensus development process, which brings together volunteers representing varied viewpoints and interests to achieve the final product. NMRA Standards documents are developed by volunteers with modeling, railroading, engineering, and industry-based expertise. Volunteers are not necessarily members of NMRA, and participate without compensation from NMRA.

NMRA does not warrant or represent the accuracy or completeness of the material contained in NMRA Standards documents, and expressly disclaims all warranties (express, implied and statutory) not included in this or any other document relating to the standard or recommended practice, including, but not limited to, the warranties of: merchantability; fitness for a particular purpose; non-infringement; and quality, accuracy, effectiveness, currency, or completeness of material. In addition, NMRA disclaims any and all conditions relating to results and workmanlike effort. In addition, NMRA does not warrant or represent that the use of the material contained in NMRA Standards documents is free from patent infringement. NMRA Standards documents are supplied "AS IS" and "WITH ALL FAULTS."

Use of NMRA Standards documents is wholly voluntary. The existence of an NMRA Standard or Recommended Practice does not imply that there are no other ways to produce, test, measure, purchase, market, or provide other goods and services related to the scope of the NMRA Standards documents. Furthermore, the viewpoint expressed at the time that NMRA approves or issues a Standard or Recommended Practice is subject to change brought about through developments in the state of the art and comments received from users of NMRA Standards documents.

In publishing and making its standards available, NMRA is not suggesting or rendering professional or other services for, or on behalf of, any person or entity, nor is NMRA undertaking to perform any duty owed by any other person or entity to another. Any person utilizing any NMRA Standards document, should rely upon their own independent judgment in the exercise of reasonable care in any given circumstances or, as appropriate, seek the advice of a competent professional in determining the appropriateness of a given NMRA Standards documents.

IN NO EVENT SHALL NMRA BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO: THE NEED TO PROCURE SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE PUBLICATION, USE OF, OR RELIANCE UPON ANY STANDARD OR RECOMMENDED PRACTICE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE AND REGARDLESS OF WHETHER SUCH DAMAGE WAS FORESEEABLE.

Translations

NMRA's development of NMRA Standards documents involves the review of documents in English only. In the event that an NMRA Standards document is translated, only the English version published by NMRA is the approved NMRA Standards document.

Official Statements

A statement, written or oral, that is not processed in accordance with NMRA policies for distribution of NMRA communications, or approved by the Board of Directors, an officer or committee chairperson, shall not be considered or inferred to be the official position of NMRA or any of its committees and shall not be considered to be, nor be relied upon as, a formal position of NMRA.

Comments on Standards

Comments for revision of NMRA Standards documents are welcome from any interested party, regardless of membership. However, **NMRA does not provide interpretations, consulting information, or advice pertaining to NMRA Standards documents.**

Suggestions for changes in documents should be in the form of a proposed change of text, together with appropriate supporting comments. Since NMRA standards represent a consensus of concerned interests, it is important that any responses to comments and questions also receive the concurrence of a balance of interests. For this reason, NMRA, its departments, Working Groups or committees cannot provide an instant response to comments, or questions except in those cases where the matter has previously been addressed. For the same reason, NMRA does not respond to interpretation requests. Any person who would like to participate in evaluating comments or in revisions to NMRA Standards documents may request participation in the relevant NMRA working group.

Laws & Regulations


Users of NMRA Standards documents should consult all applicable laws and regulations. Compliance with the provisions of any NMRA Standards document does not constitute compliance to any applicable regulatory requirements. Implementers of the standard are responsible for observing or referring to the applicable regulatory requirements. NMRA does not, by the publication of NMRA Standards documents, intend to urge action that is not in compliance with applicable laws, and NMRA Standards documents may not be construed as doing so.

Copyrights

NMRA Standards documents are copyrighted by NMRA under US and international copyright laws. They are made available by NMRA and are adopted for a wide variety of both public and private uses. These include both use, by reference, in laws and regulations, and use in private self-regulation, standardization, and the promotion of modeling, structural and engineering practices and methods. By making NMRA Standards documents available for use and adoption by public authorities and private users, NMRA does not waive any rights in copyright to the NMRA Standards documents.

IMPORTANT NOTICE

NMRA Standards documents do not guarantee or ensure safety, security, health, or environmental protection, or ensure against interference with or from other systems, devices or networks. NMRA Standards documents development activities consider research and information presented to the standards development group in developing any safety recommendations. Other information about safety practices, changes in technology or technology implementation, or impact by peripheral systems also may be pertinent to safety considerations during implementation of the standard. Implementers and users of NMRA Standards documents are responsible for determining and complying with all appropriate safety, security, environmental, health, and interference protection practices and all applicable laws and regulations.



OpenLCB Technical Note

Simple Node Information Protocol

July 22, 2024

Adopted

1 Introduction

One of the important innovations of OpenLCB compared to other layout control systems is the ability for end-users to centrally browse, manage and configure nodes connected to the OpenLCB network. This management and configuration is usually achieved by a computer software with a Graphical User Interface, such as JMRI.

These user interfaces need three major features: node discovery, node presentation in a user-friendly way (in other words *read access*), and node configuration (in other words *write access*).

The protocols defined for these purposes have to be simple enough for all nodes to implement, and robust and scalable enough to work from the smallest networks to the largest layouts or module meets.

The Simple Node Information Protocol is a light-weight and scalable protocol designed for the middle use-case: to allow a network browser tool to read a small set of node identification information (such as the manufacturer and model, as well as user-specified name and description) efficiently. This information can then be used to present the node in a UI and allow the user to recognize the node in a possibly long list or table of nodes, and to select it as a target for possible interactions.

OpenLCB provides a very versatile protocol standard for communicating between nodes, which allows very generic transfers to take place between nodes, including accessing read-only or read-write memory areas of the target node, for example via the Memory Configuration Protocol. This protocol can be parametrized to transmit the same information as the Simple Node Information Protocol.

The Simple Node Information Protocol on the other hand has been designed with simplicity and efficiency in mind. It is able to provide only a fixed and very small set of information. This makes it easy to supply data to fixed UI elements, or even to very limited UI (e.g. a single-line alphanumeric display as it could be found on a throttle). Also the implementation state flows are very simple, and allow this protocol to be implemented on the smallest (and cheapest) nodes with very minimal engineering effort. This is an important consideration, because manufacturers are naturally interested in driving the costs down, but the implementation should still support the node browsing use-cases. Another important consideration is efficiency: how much network traffic is required to retrieve basic information per node. This becomes important for very large networks, where the mere enumeration of the available nodes could lead to network congestion.

Larger nodes are therefore encouraged to still implement this protocol to allow better handling of the entire network by network management tools.

35 For a Node to implement the Simple Node Information Protocol (as a target node), the following needs to be done:

- In the Protocol Support Reply message, indicate the availability of the Simple Node Information Protocol (SNIP) protocol.
- Listen for incoming Simple Node Information Requests.
- 40 • Reply to each of these requests by sending back one or more addressed messages to the source Node, containing the fixed payload. Ensure the framing bits are appropriately set when using CAN.
- When a write operation (via any Protocol) changes the user-specified information, reset the node or transition to Uninitialized state, and transition back to Initialized state by sending Node Initialization Complete message.

45 **2 Annotations to the Standard**

2.1 Introduction

Note that this section of the Standard is informative, not normative.

2.2 Intended Use

Note that this section of the Standard is informative, not normative.

50 **2.3 Reference and Context**

The OpenLCB Message Network Standard defines the CAN frame encoding of addressed OpenLCB messages. The reply generated by the Simple Node Information Protocol is an addressed message, and thus is governed by this formatting requirement. While CAN-bus messages can only transmit 8 bytes of data at a time, the response payload to send can be up to 253 bytes long. This requires the use of
55 multiple CAN frames. The OpenLCB Message Network Standard defines two bits in the first data byte which allows sending multiple frames that constitute a single addressed message.

The Protocol Support Inquiry and Protocol Support Reply messages allow a node (for example the network browser node) to check what protocols a particular destination node supports. If a node implements the Simple Node Information Protocol, the network browser will know from the Protocol
60 Support Reply.

The OpenLCB protocol stack provides a number of methods, some very flexible, to exchange data between participating nodes. One of these methods is laid out to represent the same information as provided by the Simple Node Information Protocol, albeit in a more general way. The OpenLCB Datagram Transport Protocol can be used to send medium-sized messages between nodes, the
65 OpenLCB Memory Configuration Protocol can be used on top of this to provide read-only or read-write access to blocks of memory in the target node. The exported blocks are organized as so-called Memory Spaces, and two of these spaces are reserved to provide the same information as the Simple Node Information Protocol. The space 0xFC gives access to the read-only manufacturer-specific data, while the space 0xFB gives access to the read-write user-modifiable data. These two data blocks
70 together give the information returned by the Simple Node Information Protocol. The format is described in the OpenLCB Configuration Description Information Standard.

For a Node to implement the *alternative* retrieval method for this information, the following steps have to be made:

- 75 • In the Protocol Support Reply, indicate the availability of the Abbreviated Configuration Definition Information (ACDI) protocol.
- Implement the OpenLCB Datagram Transport Protocol and the OpenLCB Memory Access Configuration Protocol.
- Implement a read-only memory space 0xFC and export the manufacturer-specific identification information in a fixed format.
- 80 • Implement a read-write memory space 0xFB and export the user-specific identification information in a fixed format.
- If the Configuration Definition Information (CDI) memory space is implemented, then add the `<acdi ... />` tag to the XML description of the node configuration.

For details, see the above quoted standards.

85 **2.4 Message Format**

This section of the Standard is normative, and describes the OpenLCB messages in a generic form (transport-agnostic). For examples of how these messages would look on CAN, see [Section 2.7](#).

2.5 Payload Format

2.5.1 Format

- 90 The total size of the reply can be up to 253 bytes. A typical reply will be much shorter, because only the actual length of the strings have to be transported as opposed to the maximum length (like for example in the case of the Memory Configuration Protocol).

2.5.2 Versioning

- 95 The first version of this standard used the version identifiers 0x01 and 0x01 in the two sections, respectively. In order to allow future extensions of the data transmitted, we must ensure that nodes parsing the response according to the current version of the standard will be able to understand a response generated by a Node according to a future version of the Standard.

- 100 The solution is that the version identifier byte will double as a count of the null-terminated strings in the particular data segment, hence the current standard defines 0x04 and 0x02 as valid version numbers, because the first section has four fields and the second section has two fields. Thus a receiving node can count the number of zero bytes in the received message to find the separator between the sections and identify the known fields.

Nodes must accept data with a higher version number than in this standard, by ignoring the additional strings, as described above.

105 **2.6 Interactions**

When a node gets a Simple Node Information Request, if possible it shall reply with one or more Simple Node Information Reply messages containing the node's information. If it's not able to process the request, (for example because it is already servicing another request), it shall send an Optional

110 Interaction Rejected with an appropriate error code. It's recommended that the rejection message have the temporary-error bit set, so that the node sending the original request will retry it. The suggested error code is 0x2020, Buffer unavailable. See the OpenLCB Message Network Standard for the Optional Interaction Rejected message and for the error code format.

115 The first version of this Standard predates the ability to send long addressed messages over CAN-bus. There are Nodes produced which will send the reply in as many messages as the number of CAN frames necessary to cover the length, without using the framing bits. Therefore network browsers using the Simple Node Information Protocol must be prepared to receive the response in multiple messages. If there is no framing bits to signal the end of transmission, the receiving node may count the zero bytes to realize when the response transmission is finished.

120 While the Simple Node Information Protocol is a lightweight protocol, enumerating a large set of nodes will still put a considerable traffic onto the bus. Therefore it is strongly recommended for network browser nodes to cache the returned information in their memory to avoid unnecessary traffic in repeatedly requesting it. The Standard specifies that this cache must be invalidated when the Node Initialization Complete message is received from a particular node, and conversely, if the Node's information payload changes, it must re-enter Initialized state to notify any other node that may have
125 cached the information to refresh it.

2.7 CAN Adaptation

For reference, we describe the messages defined in this standard as formatted for the CAN transport protocol. This section is not normative; should there be a conflict between this document and the Standards, then the Standards take precedence.

Name	CAN header	Data Content
Simple Node Information Request	0x19DE,8sss	fddd

130 Where 'sss' is the 12-bit CAN Alias of the source Node (the node sending the request), 'ddd' is the 12-bit CAN Alias of the target Node (being queried), and 'f' is typically 0.

Name	CAN header	Data Content
Simple Node Information Reply	0x19A0,8sss	fddd, up to 6 payload bytes

Where 'sss' is the 12-bit CAN Alias of the target Node (sending the reply), 'ddd' is the 12-bit CAN Alias of the source Node (the one that has sent the query), and 'f' is as follows:

- 135
- f = 0x0 for replies that fit into a single CAN frame
 - f = 0x1 for the first frame of a reply with more than one frame
 - f = 0x2 for the last frame of a reply with more than one frame
 - f = 0x3 for each middle frames of a reply with more than two frames

140 Note is made again that previous node implementations may send the SNIP replies as multiple independent messages, where f is zero, and receiving nodes should accept these. The end of the SNIP can be determined by counting the fields as determined by the version numbers.

Table of Contents

1 Introduction.....1

2 Annotations to the Standard.....2

2.1 Introduction.....2

2.2 Intended Use.....2

2.3 Reference and Context.....2

2.4 Message Format.....3

2.5 Payload Format.....3

2.5.1 Format.....3

2.5.2 Versioning.....3

2.6 Interactions.....3

2.7 CAN Adaptation.....4